

## **Method of Procedure (MoP) to implement the “End-to-End Telecom Pipeline” in Python + SQLite.**

You will:

- Load raw usage data from telecom\_raw.csv (e.g. extracting from data lake)
- Clean & transform it, compute total usage per customer and flag heavy-user (ETL pipeline)
  - Aggregate total usage per customer
  - Flag heavy users (>50GB OR >500 minutes)
- Load the cleaned data into a SQLite “mini data-warehouse” with an index.
- Create role-based views for junior analysts vs managers.
- Automate the pipeline to run for 20 seconds for 3 times, with simple logging.

### 1. Environment setup

Open Anaconda & Jupyter -> Start Anaconda Navigator -> Click Jupyter Notebook.  
In the browser tab, create a new folder on the Desktop/telecom\_pipeline.  
Click new notebook pipeline.ipynb.

Install required libraries (one-time)

```
pip install pandas schedule
```

### 2. Prepare raw data file (telecom\_raw.csv)

For class practice we'll create a small dummy CSV directly from Python.

```
import pandas as pd
raw_data = [
    # customer_id, call_minutes, data_gb, timestamp, region
    [1001, 120, 10.5, "2025-09-25 10:30", "Delhi"],
    [1002, 450, 55.0, "2025-09-25 11:00", "Mumbai"],
    [1003, 200, 20.0, "25-09-2025 09:15", "Chennai"], # different date format
    [1004, 600, 70.0, "2025/09/25 08:00", "Delhi"], # different date format
    [1005, 50, 5.0, "2025-09-25 12:00", "Kolkata"],
    [1005, 50, 5.0, "2025-09-25 12:00", "Kolkata"], # duplicate
    [1006, None, 60.0, "2025-09-25 13:00", "Hyderabad"],# missing call_minutes
    [1007, 510, None, "2025-09-25 14:00", "Delhi"], # missing data_gb
]

columns = ["customer_id", "call_minutes", "data_gb", "timestamp", "region"]
df_raw = pd.DataFrame(raw_data, columns=columns)

df_raw.to_csv("telecom_raw.csv", index=False)
print("Saved telecom_raw.csv")
df_raw
```

### 3. Common imports and database file

```
import sqlite3
import logging
import schedule
import time
```

```

import os
import pandas as pd
#Set database path and configure logging

DB_PATH = "telecom_warehouse.db"

# Basic logging configuration
logging.basicConfig(
    filename="telecom_pipeline.log",
    level=logging.INFO,
    format"%(asctime)s [%(levelname)s] %(message)s",
)
print("Database path:", DB_PATH)

```

#### 4. ETL Step - Extract, Transform & Clean

```

# Extract
def extract_raw_data(csv_path="telecom_raw.csv"):
    """
    Step 1: Extract raw data from CSV file.
    Returns a pandas DataFrame.
    """

    if not os.path.exists(csv_path):
        raise FileNotFoundError(f"{csv_path} not found in current folder.")

    df = pd.read_csv(csv_path)
    logging.info("Extracted %d rows from %s", len(df), csv_path)
    return df

```

#### **# Function to transform & clean**

```

import pandas as pd
import logging

# Function to transform & clean
def transform_and_clean(df_raw):
    """
    Step 2: Clean and transform the data.
    - Remove exact duplicate rows
    - Standardize date format
    - Handle missing values (impute)
    - Aggregate total usage per customer
    - Flag heavy users (>50GB OR >500 minutes)
    """

    df = df_raw.copy()

    # 1) Remove duplicates
    before = len(df)
    df = df.drop_duplicates()
    logging.info("Removed %d duplicate rows", before - len(df))

```

```

# 2) Standardize timestamp to datetime
df["timestamp"] = pd.to_datetime(df["timestamp"], errors="coerce")
bad_dates = df["timestamp"].isna().sum()
if bad_dates > 0:
    logging.warning("Found %d bad timestamps; filling with default date", bad_dates)
    df["timestamp"] = df["timestamp"].fillna(pd.Timestamp("2025-09-25 00:00"))

# 3) Handle missing numeric values: fill with column mean
for col in ["call_minutes", "data_gb"]:
    if df[col].isna().sum() > 0:
        mean_val = df[col].mean()
        df[col] = df[col].fillna(mean_val)
        logging.info("Filled missing values in %s with mean=%,.2f", col, mean_val)

# 4) Aggregate total per customer
agg = (
    df.groupby(["customer_id", "region"])
    .agg(
        total_call_minutes=("call_minutes", "sum"),
        total_data_gb=("data_gb", "sum"),
        last_activity=("timestamp", "max"),
    )
    .reset_index()
)

# 5) Flag heavy users
agg["heavy_user"] = (agg["total_data_gb"] > 50) | (agg["total_call_minutes"] > 500)

logging.info("Transformed data to %d aggregated customer rows", len(agg))
return agg

# Use the function, save cleaned data to a NEW CSV, and print it ---
df_raw = extract_raw_data("telecom_raw.csv") # raw CSV is only read, not modified
df_clean = transform_and_clean(df_raw)

# Save cleaned/aggregated data to a separate CSV
clean_csv_path = "telecom_cleaned.csv"
df_clean.to_csv(clean_csv_path, index=False)

print("Cleaned dataset (first few rows):")
print(df_clean.head())

print(f"\nCleaned data has been saved to: {clean_csv_path}")

```

You should see one row per customer with total\_call\_minutes, total\_data\_gb, last\_activity, and heavy\_user.

## 5. Load to SQLite “warehouse”

We now store the cleaned data in a SQLite DB and create an index on customer\_id.

```
def load_to_warehouse(df_clean, db_path=DB_PATH):
    """
    Load cleaned data into SQLite as 'usage_summary' table.
    Creates an index on customer_id for fast lookups.
    """
    conn = sqlite3.connect(db_path)
    cur = conn.cursor()

    # Write DataFrame to table (replace for each run)
    df_clean.to_sql("usage_summary", conn, if_exists="replace", index=False)

    # Create index on customer_id
    cur.execute("CREATE INDEX IF NOT EXISTS idx_usage_summary_cust ON
usage_summary(customer_id);")

    conn.commit()
    conn.close()
    logging.info("Loaded %d rows into usage_summary table", len(df_clean))

load_to_warehouse(df_clean)

# Quick check: read from DB and show
conn = sqlite3.connect(DB_PATH)
check_df = pd.read_sql_query("SELECT * FROM usage_summary;", conn)
conn.close()

print("Data inside SQLite usage_summary table:")
check_df
```

You should see the same cleaned rows coming from SQLite.

## 6. Apply Governance (role-based views)

We'll create two SQL views:

vw\_usage\_junior – hides names/phones  
vw\_usage\_manager – full usage table.

For this mini-project we'll simulate PII fields (customer\_name, phone) inside SQLite and show how views would mask them.

#Enrich table with fake PII fields

```
def enrich_with_pii(db_path=DB_PATH):
    """
    Adds fake PII columns to usage_summary to demonstrate role-based masking.
    """
    conn = sqlite3.connect(db_path)
    cur = conn.cursor()
```

```

# Add columns if they don't exist
cur.execute("PRAGMA table_info(usage_summary);")
cols = [row[1] for row in cur.fetchall()]

if "customer_name" not in cols:
    cur.execute("ALTER TABLE usage_summary ADD COLUMN customer_name TEXT;")
if "phone" not in cols:
    cur.execute("ALTER TABLE usage_summary ADD COLUMN phone TEXT;")

# Simple mapping for demo
name_map = {
    1001: "Asha Mehta",
    1002: "Ravi Kumar",
    1003: "Sneha Rao",
    1004: "Manoj Singh",
    1005: "Divya Jain",
    1006: "Rahul Roy",
    1007: "Neha Gupta",
}
}

# Update rows with names and fake phone numbers
for cust_id, name in name_map.items():
    cur.execute(
        """
        UPDATE usage_summary
        SET customer_name = ?, phone = ?
        WHERE customer_id = ?;
        """,
        (name, f"9{cust_id}000000", cust_id)
    )

conn.commit()
conn.close()
logging.info("Enriched usage_summary with fake PII fields.")

enrich_with_pii()

```

### **Next Step – Create role-based SQL views**

```

def create_role_based_views(db_path=DB_PATH):
    """
    Create role-based views:
    - vw_usage_junior: masked PII
    - vw_usage_manager: full access
    """

    conn = sqlite3.connect(db_path)
    cur = conn.cursor()

    # Junior view: NO real name/phone, only masked info
    cur.execute("""
        CREATE VIEW IF NOT EXISTS vw_usage_junior AS
    """)

```

```

SELECT
    customer_id,
    region,
    total_call_minutes,
    total_data_gb,
    last_activity,
    heavy_user,
    'ANON' AS customer_name,
    'XXXXXX' AS phone
FROM usage_summary;
""")

# Manager view: full details
cur.execute("""
CREATE VIEW IF NOT EXISTS vw_usage_manager AS
SELECT
    customer_id,
    region,
    total_call_minutes,
    total_data_gb,
    last_activity,
    heavy_user,
    customer_name,
    phone
FROM usage_summary;
""")

conn.commit()
conn.close()
logging.info("Created role-based views for junior and manager roles.")

create_role_based_views()

```

### **Next Step – Helper function to read data for a given role**

```

def get_usage_for_role(role, db_path=DB_PATH, limit=10):
    """
    Utility: return a DataFrame for the requested role ('junior' or 'manager').
    """

    view_map = {
        "junior": "vw_usage_junior",
        "manager": "vw_usage_manager",
    }
    view_name = view_map.get(role.lower())
    if view_name is None:
        raise ValueError("Role must be 'junior' or 'manager'")

    conn = sqlite3.connect(db_path)
    df = pd.read_sql_query(f"SELECT * FROM {view_name} LIMIT {limit};", conn)
    conn.close()
    return df

```

```
print("Junior analyst view:")
display(get_usage_for_role("junior"))

print("\nManager view:")
display(get_usage_for_role("manager"))
```

You should see that the junior view has customer\_name = 'ANON' and phone = 'XXXXXX', while the manager view shows real names and phone numbers.

## 7. Automate the pipeline (every 15 seconds)

Let's wrap everything into a single function run\_pipeline() and then schedule it. Define the full pipeline function:

```
def run_pipeline():
    """
    Runs the full pipeline once:
    1) Extract raw data from CSV
    2) Transform & clean
    3) Load to warehouse (SQLite)
    4) Enrich with PII and create role-based views
    """

    try:
        logging.info("Starting ETL pipeline run...")

        # 1 & 2: Extract + Transform
        df_raw = extract_raw_data("telecom_raw.csv")
        df_clean = transform_and_clean(df_raw)

        # 3: Load
        load_to_warehouse(df_clean, DB_PATH)

        # 4: Governance
        enrich_with_pii(DB_PATH)
        create_role_based_views(DB_PATH)

        logging.info("ETL pipeline run completed successfully.")

    except Exception as e:
        logging.exception("ETL pipeline run FAILED: %s", e)
        print("Pipeline failed, see log file for details.")
```

### Next Step – Test running the pipeline once

```
run_pipeline()

print("Manager view after pipeline run:")
get_usage_for_role("manager")
```

If this works and the table/views look correct, you're ready to automate.

## 8. Scheduling the pipeline (simulated every 20 seconds)

For classroom demos, let it run for 3 iterations then stop it. Schedule job every 20 seconds  
import time, schedule

```
schedule.clear()  
schedule.every(20).seconds.do(run_pipeline)
```

```
runs = 3 # run 3 times then stop  
for _ in range(runs):  
    schedule.run_pending()  
    time.sleep(10) # wait for the next tick  
  
print("✅ Done. Scheduler exited after", runs, "runs.")
```

Watch the notebook; every ~15 seconds, a new log entry will be added in telecom\_pipeline.log and the SQLite file telecom\_warehouse.db will be updated.

## 9. Connecting SQLite with Power BI for Visualization

Follow the steps below to connect your telecom\_warehouse.db SQLite database to Power BI and build visual dashboards using the cleaned dataset and governance views.

### Step 1 — Install SQLite ODBC Driver

Power BI does not natively support SQLite, so an ODBC connector is required.

- Download a **64-bit SQLite ODBC driver** (matching Power BI Desktop architecture).
- Install using default settings.
- Restart Power BI if already open.

### Step 2 — Open Power BI Desktop

- Launch **Power BI Desktop**
- Select **Get Data → More...**

### Step 3 — Choose ODBC Source

- Search for **ODBC**
- Click **Connect**

### Step 4 — Select or Create SQLite Data Source

If a SQLite DSN already exists, select it.

If not:

1. Click **DSN Manager**
2. Create a new DSN
3. Select: **SQLite ODBC Driver**
4. Browse and point it to:

telecom\_warehouse.db

Click **OK**.

### Step 5 — Select the Table or View

Power BI will list all database objects including:

Type Name	Access Level
Table usage_summary	Full cleaned dataset
View vw_usage_junior	Masked data (no PII)
View vw_usage_manager	Full PII access
Choose based on reporting/audience needs.	
Then click:	
→ Load	

### Step 6 — Build Visuals

Once data loads, you can create visualizations such as:

- **Bar Chart:** Top heavy users
- **Stacked Chart:** Data usage by region
- **KPI Cards:** Total call minutes, total data consumption
- **Filters:** Region, heavy user flag, time period

10. What students should submit

1. telecom\_raw.csv
2. telecom\_warehouse.db
3. Notebook telecom\_pipeline.ipynb
4. Screenshot of:
  - Junior vs manager views, or
  - Power BI dashboard